

DataSet: Pima Indians Diabetes Dataset

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

```
1 from sklearn.model_selection import train_test_split
2 from sklearn.preprocessing import LabelEncoder, StandardScaler
3 from sklearn.metrics import classification_report, confusion_matrix
4 from sklearn.metrics import ConfusionMatrixDisplay
5 from sklearn.model_selection import learning_curve
```

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.neighbors import KNeighborsClassifier
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.ensemble import RandomForestClassifier
5 from sklearn.svm import SVC
```

```
1 df = pd.read_csv("/content/drive/MyDrive/Dataset/Pima Indians Dia
```

```
1 display(df.head())
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Dia
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
1 print(df.shape)
```

```
(768, 9)
```

```
1 print(df.columns.tolist())
```

```
['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin'
```

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```
1 display(df.describe())
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

```
1 print((df == 0).sum())
```

Pregnancies	111
Glucose	5
BloodPressure	35
SkinThickness	227
Insulin	374
BMI	11
DiabetesPedigreeFunction	0
Age	0
Outcome	500

dtype: int64

```
1 zero_to_nan_cols = ["Glucose", "BloodPressure", "SkinThickness",  
2
```

```
1 df[zero_to_nan_cols] = df[zero_to_nan_cols].replace(0, np.nan)
```

```
1 print(df.isnull().sum())
```

Pregnancies	0
Glucose	5

```
BloodPressure      35
SkinThickness     227
Insulin          374
BMI              11
DiabetesPedigreeFunction 0
Age              0
Outcome          0
dtype: int64
```

```
1 df.fillna(df.median(), inplace=True)
```

```
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction 0
Age              0
Outcome          0
dtype: int64
```

```
1 before_rows = df.shape[0]
2 df.drop_duplicates(inplace=True)
3 after_rows = df.shape[0]
```

```
1 print(f"\nRows Before Removing Duplicates: {before_rows}")
2 print(f"Rows After Removing Duplicates: {after_rows}")
```

```
Rows Before Removing Duplicates: 768
Rows After Removing Duplicates: 768
```

```
1 print(df["Outcome"].unique())
```

```
[1 0]
```

```
1 le = LabelEncoder()
2 df["Outcome"] = le.fit_transform(df["Outcome"])
```

```
1
2 print("\nUnique values in Outcome after encoding:")
3 print(df["Outcome"].unique())
```

```
Unique values in Outcome after encoding:
[1 0]
```

```
1 X = df.drop("Outcome", axis=1)
2 y = df["Outcome"]
```

```
1 print("\nFeature Matrix (X) Shape:", X.shape)
2 print("Target Vector (y) Shape:", y.shape)
```

Feature Matrix (X) Shape: (768, 8)
Target Vector (y) Shape: (768,)

```
1 scaler = StandardScaler()
2 X_scaled = scaler.fit_transform(X)
3
4 print("\nFeature Scaling Completed.")
```

Feature Scaling Completed.


```
1 X_train, X_test, y_train, y_test = train_test_split(
2     X_scaled, y, test_size=0.20, random_state=42
3 )
```

```
1 print("\nTraining Set Shape:", X_train.shape)
2 print("Testing Set Shape:", X_test.shape)
```

Training Set Shape: (614, 8)
Testing Set Shape: (154, 8)


```
1 print("\nGenerating Histograms...")
2 df.hist(figsize=(14, 10), bins=20)
3 plt.tight_layout()
4 plt.show()
```

Generating Histograms...



```
1 print("\nCorrelation Heatmap:")
2 plt.figure(figsize=(10, 6))
3 sns.heatmap(df.corr(), annot=True, cmap="coolwarm")
4 plt.show()
```

Correlation Heatmap:




```


1 print("\nScatter Plot: Glucose vs BMI")
2 plt.figure(figsize=(6, 4))
3 sns.scatterplot(x=df["Glucose"], y=df["BMI"], hue=df["Outcome"])
4 plt.show()

```

Scatter Plot: Glucose vs BMI



```
1 plt.figure(figsize=(6, 5))
2 plt.plot(fpr, tpr, label="Logistic Regression")
3 plt.plot([0, 1], [0, 1], linestyle="--")
4 plt.xlabel("False Positive Rate")
5 plt.ylabel("True Positive Rate")
6 plt.title("ROC Curve")
7 plt.legend()
8 plt.show()
```




```
1 print("\nGenerating Learning Curve for SVM...")  
2  
3 train_sizes, train_scores, test_scores = learning_curve(  
4     SVC(), X_scaled, y, cv=5, scoring='accuracy'  
5 )
```

Generating Learning Curve for SVM...

```
1 train_mean = train_scores.mean(axis=1)  
2 test_mean = test_scores.mean(axis=1)
```


```
1 plt.figure(figsize=(7, 5))  
2 plt.plot(train_sizes, train_mean, label="Training Score")  
3 plt.plot(train_sizes, test_mean, label="Validation Score")  
4 plt.xlabel("Training Set Size")  
5 plt.ylabel("Accuracy")  
6 plt.title("Learning Curve (SVM)")  
7 plt.legend()  
8 plt.show()
```



```
1 log_model = LogisticRegression(max_iter=300)
2 log_model.fit(X_train, y_train)
3
4 y_pred_log = log_model.predict(X_test)
5
6 print("Confusion Matrix - Logistic Regression:")
7 cm_log = confusion_matrix(y_test, y_pred_log)
8 print(cm_log)
9
10 disp = ConfusionMatrixDisplay(confusion_matrix=cm_log)
11 disp.plot(cmap="Blues")
12 plt.title("Confusion Matrix - Logistic Regression")
13 plt.show()
14
15 print("\nClassification Report - Logistic Regression:")
16 print(classification_report(y_test, y_pred_log))
17
```

Confusion Matrix - Logistic Regression:

```
[[82 17]
 [21 34]]
```



Classification Report - Logistic Regression:

	precision	recall	f1-score	support
0	0.80	0.83	0.81	99
1	0.67	0.62	0.64	55
accuracy			0.75	154
macro avg	0.73	0.72	0.73	154
weighted avg	0.75	0.75	0.75	154


```

1 knn_model = KNeighborsClassifier()
2 knn_model.fit(X_train, y_train)
3
4 y_pred_knn = knn_model.predict(X_test)
5
6 print("Confusion Matrix - KNN:")
7 cm_knn = confusion_matrix(y_test, y_pred_knn)
8 print(cm_knn)
9
10 disp = ConfusionMatrixDisplay(confusion_matrix=cm_knn)
11 disp.plot(cmap="Blues")
12 plt.title("Confusion Matrix - KNN")
13 plt.show()
14
15 print("\nClassification Report - KNN:")
```

```
16 print(classification_report(y_test, y_pred_knn))
17
```

Confusion Matrix - KNN:

```
[[75 24]
 [17 38]]
```




Classification Report - KNN:

	precision	recall	f1-score	support
0	0.82	0.76	0.79	99
1	0.61	0.69	0.65	55
accuracy			0.73	154
macro avg	0.71	0.72	0.72	154
weighted avg	0.74	0.73	0.74	154

```
1 dt_model = DecisionTreeClassifier()
2 dt_model.fit(X_train, y_train)
3
4 y_pred_dt = dt_model.predict(X_test)
5
6 print("Confusion Matrix - Decision Tree:")
7 cm_dt = confusion_matrix(y_test, y_pred_dt)
8 print(cm_dt)
9
10 disp = ConfusionMatrixDisplay(confusion_matrix=cm_dt)
11 disp.plot(cmap="Greens")
12 plt.title("Confusion Matrix - Decision Tree")
13 plt.show()
```

```
14  
15 print("\nClassification Report - Decision Tree:")  
16 print(classification_report(y_test, y_pred_dt))
```

Confusion Matrix - Decision Tree:
[[75 24]
 [23 32]]



Classification Report - Decision Tree:


	precision	recall	f1-score	support
0	0.77	0.76	0.76	99
1	0.57	0.58	0.58	55
accuracy			0.69	154
macro avg	0.67	0.67	0.67	154
weighted avg	0.70	0.69	0.70	154

```
1  
2 rf_model = RandomForestClassifier()  
3 rf_model.fit(X_train, y_train)  
4  
5 y_pred_rf = rf_model.predict(X_test)  
6  
7 print("Confusion Matrix - Random Forest:")  
8 cm_rf = confusion_matrix(y_test, y_pred_rf)  
9 print(cm_rf)  
10  
11 disp = ConfusionMatrixDisplay(confusion_matrix=cm_rf)  
12 disp.plot(cmap="Purples")
```

```
13 plt.title("Confusion Matrix - Random Forest")
14 plt.show()
15
16 print("\nClassification Report - Random Forest:")
```

Confusion Matrix - Random Forest:

```
[[80 19]
 [17 38]]
```



Classification Report - Random Forest:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.82	0.81	0.82	99
1	0.67	0.69	0.68	55

accuracy		0.77	154
----------	--	------	-----